

Advanced Algorithms

Lecture Notes

Luca Simonetti

2025–2026

Contents

1	String Matching	5
1.1	Introduction & Exact Pattern Matching	5
1.1.1	Definition	5
1.1.2	Naive approach	6
1.1.3	Knuth-Morris-Pratt (KMP) algorithm	8
1.1.4	Z-Algorithm	17
1.1.5	Strong borders and real-time KMP	22
1.1.6	A Comprehensive Trace: KMP and Z-Algorithm	24
1.1.7	Boyer-Moore Algorithm	26
2	Randomized Algorithms	29
3	Data Compression	31

Chapter 1

String Matching

1.1 Introduction & Exact Pattern Matching

1.1.1 Definition

The exact pattern matching problem is defined as follows:

Problem 1.1.1 (Exact pattern matching). Given an alphabet Σ , and two strings $T, P \in \Sigma^*$, for which $|T| = n$ and $|P| = m$, find all occurrences of P in T , i.e. compute $\{i \mid T[i, i + m - 1] = P\}$.

Notation

Throughout these notes, Σ denotes a finite alphabet of constant size ($|\Sigma| \in \mathcal{O}(1)$). For a string S of length $|S|$:

- $S[i]$ denotes the i -th character of S (1-based indexing).
- $S[i, j]$ denotes the substring $S[i]S[i + 1] \cdots S[j]$ (empty string if $i > j$).
- $S[j] = S[1, j]$ is the prefix of length j ; $S[i,] = S[i, |S|]$ is the suffix starting at i .

A *proper* prefix (resp. suffix) of S is any prefix (resp. suffix) strictly shorter than S itself.

Example 1.1.2. Let $\Sigma = \{a, b\}$, $T = \text{ababbab}$ and $P = \text{abb}$. The only occurrence of P in T is at position 3.

1.1.2 Naive approach




Figure 1.1: Naive exact pattern matching of $P = \text{abb}$ in $T = \text{ababbab}$ (1-indexed). Each row corresponds to one shift s ; the pattern P is aligned starting at $T[s]$. **Green** = character matched, **red** = first mismatch (algorithm skips to next shift), gray = not compared.

What the Figure 1.1 illustrates is one basic naive approach. In this approach we nest two loops: the outer loop iterates over all possible "first positions" s in T where the pattern P could potentially match, and the inner loop checks character by character if P matches T starting from position s . For example the first row of the figure corresponds to starting the search from index $s = 1$ of T . Given s we then loop one by one each character of P , moving along T as well, and check if all the characters match. In this case we find a mismatch at the third character of P (which is b , whereas the corresponding character in T is a). When this happens, we break this inner loop and move to the next shift $s = 2$. If the length of the pattern was m and the length of the text was n , the outer loop could have up to $\mathcal{O}(n)$ iterations, and each inner iteration could take up to $\mathcal{O}(m)$ time, leading to a worst-case time complexity of $\mathcal{O}(nm)$ for this naive approach.




Figure 1.2: Worst-case input for the naive algorithm: $T = \text{aaaaaaa}$ ($n = 7$) and $P = \text{aaab}$ ($m = 4$). Every shift s performs $m - 1$ successful comparisons before the inevitable mismatch on $P[m] = \text{b}$, giving $(n - m + 1)(m - 1) = \Theta(nm)$ total comparisons.

In figure 1.2 we see an example of a worst-case input for the naive algorithm. In this case the text T consists of n a characters, and the pattern P consists of $m - 1$ a characters followed by a b . There are $n - m + 1$ possible shifts s where

the pattern could potentially match, and for each of these shifts the algorithm will successfully match the first $m - 1$ characters of P . In the example of the figure we have $n = 7$ and $m = 4$, so there are $7 - 4 + 1 = 4$ shifts to check, and for each shift we make m checks, of which $m - 1 = 3$ are successful matches before we encounter the mismatch at the last character of P . This results in a total of $(n - m + 1)(m - 1)$ comparisons, which is $\Theta(nm)$ in the worst case.

Algorithm 1: Naive exact pattern matching

Input: Text $T[1 \dots n]$, Pattern $P[1 \dots m]$

Output: All positions s such that $T[s, s + m - 1] = P$

```

1 for  $s \leftarrow 1$  to  $n - m + 1$  do
2    $j \leftarrow 1$ 
3   while  $j \leq m$  and  $T[s + j - 1] = P[j]$  do
4      $j \leftarrow j + 1$ 
5   if  $j > m$  then
6     output  $s$ 

```

Correctness

To pin down what “the inner loop does” precisely, let’s give a name to the obvious quantity: for any position i in T , write

$$\text{common}(i) = \max\{k \geq 0 \mid T[i, i + k - 1] = P[1, k]\}$$

for the length of the longest common prefix of $T[i, \dots]$ and P — basically, how many characters of P would match if you aligned it at position i .

Lemma 1.1.3. *At shift s , the inner loop exits with $j = \text{common}(s) + 1$.*

Proof. Track the invariant: entering each iteration, we have $T[s, s + j - 2] = P[1, j - 1]$ (the first $j - 1$ characters already matched). The loop keeps going as long as the next character also matches; it stops the moment it doesn’t, or when j runs past m . Since $\text{common}(s)$ is exactly how many consecutive characters agree, the loop stops right at $j = \text{common}(s) + 1$. $\square$

Theorem 1.1.4. *The naive algorithm finds every occurrence of P in T , and nothing more.*

Proof. The outer loop visits every shift $s \in \{1, \dots, n - m + 1\}$, so no candidate starting position is ever skipped. By Lemma 1.1.3, after the inner loop we have $j = \text{common}(s) + 1$; the algorithm outputs s exactly when $j > m$, i.e. $\text{common}(s) \geq m$, i.e. $T[s, s + m - 1] = P$ — which is precisely the definition of an occurrence. $\square$

Lower bound

Theorem 1.1.5. *Any algorithm for exact pattern matching must make $\Omega(n + m)$ character comparisons in the worst case.*

Proof. Two independent arguments, one for each term.

- **You have to read all of P ($\Omega(m)$).** Suppose the algorithm skips position $P[q]$. An adversary flips just that character, producing a different pattern P' . The algorithm sees the same comparisons on both inputs and therefore gives the same output — which can't be correct for both P and P' .
- **You have to read all of T ($\Omega(n)$).** Suppose position $T[i]$ is never read. The adversary flips that character, potentially creating or destroying a match right around position i . Since the algorithm never looked at $T[i]$, it has no way to tell the difference.

The two bounds add up: $\Omega(m) + \Omega(n) = \Omega(n + m)$. $\square$

The naive algorithm achieves $\mathcal{O}(nm)$ in the worst case, which is far from this lower bound. KMP will close the gap, achieving $\Theta(n + m)$.

This $\Omega(n + m)$ lower bound applies to any algorithm, not just naive.

Properties

We evaluate string matching algorithms along four axes:

Online. The algorithm processes T strictly left-to-right, one character at a time, and can emit results before reading future characters of T . Useful when T arrives as a stream.

Real-time. A stronger requirement: each new character of T is fully processed in $\mathcal{O}(1)$ *worst-case* time before the next one is read.

Blind (no backtracking). The algorithm never re-reads a position of T : each character of T is accessed at most once. A blind algorithm is necessarily online; the converse is false (the naive algorithm is online but not blind).

Succinct. The algorithm uses only $\mathcal{O}(1)$ extra space beyond storing P and T (no large auxiliary data structures).

Algorithm	Online	Real-time	Blind	Succinct
Naive	✓	×	×	✓
KMP	✓	×	✓	×
KMP (sp')	✓	✓	✓	×

The naive algorithm backtracks in T : at each new shift s it re-reads characters already compared at shift $s - 1$.

The naive algorithm is **online** (it reads T left-to-right) but not **blind**: at each new shift it backtracks and re-compares characters of T that were already seen. It is also not **real-time**: a single character of T can trigger up to $\mathcal{O}(m)$ comparisons.

This naive approach basically starts over from scratch at each shift, without leveraging precious information from previous comparisons.

1.1.3 Knuth-Morris-Pratt (KMP) algorithm

Intuition

What if we could somehow "remember" the information from a previous comparison, and use it to skip some of the checks at the next shift? Let's introduce the intuition behind the Knuth-Morris-Pratt (KMP) algorithm, with an example.

Example 1.1.6. Let $P = \text{abcdeabdef}$ and suppose we are running the naive algorithm at shift $s = 3$ over some text T whose characters at positions 3–9 are abcdeab . We match $P[1 \dots 7]$ successfully, then hit a mismatch at $P[8] = \text{d}$ vs $T[10] = \text{x}$.

As shown in Figure 1.3, at this point the naive algorithm would discard everything and try shift $s = 4$. But notice that the matched portion abcdeab starts *and* ends with $\alpha = \text{ab}$: the prefix $P[1 \dots 2]$ and the suffix $P[6 \dots 7]$ of the matched region are identical. This means that when we slide P forward, the α that was at the end of the match is now perfectly aligned with the α at the start of P — we can jump directly to shift $s = 8$ and resume comparing from $P[3]$, skipping five shifts entirely.





Figure 1.3: KMP shift intuition with $P = \text{abcdeabdef}$. At shift $s = 3$, the algorithm matches $P[1 \dots 7] = \text{abcdeab}$ then hits a mismatch at $P[8]$. The matched portion ends with $\alpha = \text{ab}$, which also appears at the very start of P — so instead of shifting by 1, we jump directly to $s = 8$, re-aligning α and resuming from $P[3]$.

Basically what the example shows is that if we take the longest prefix of P that is also a suffix of the **matched** portion of P , we can shift the pattern so that this prefix is now aligned with the suffix — and we can resume the search from the character right after this prefix, since we know it must match (otherwise we would have had a mismatch earlier). In other words we have saved some work (namely, five shift indexes) by leveraging the information from the previous comparisons, instead of starting over from scratch, from the very next shift.

But is this correct? Or do we miss some potential matches by jumping directly to $s = 8$?

When we say suffix here we mean suffix of the **matched** portion of P , not of P itself. In the example, α is a suffix of $P[1 \dots 7]$ but not of the whole P .



“You miss 100% of the shots you don’t take.” — Michael Scott

Well, it turns out that this is indeed correct, and we won’t miss any matches

by doing this. The KMP algorithm is based on this very idea, and it guarantees that we will find all occurrences of P in T without missing any, while also achieving a much better time complexity than the naive approach.

Preprocessing: computing sp values

To exploit this idea, we need to precompute for every prefix $P[1 \dots i]$ its *longest proper prefix that is also a suffix*. We write this length as $sp_i(P)$.

A proper prefix/suffix excludes the string itself; $sp_i(P)$ is thus an integer.

Definition 1.1.7 (sp values). Let P be a pattern of length m . For each $i \in \{1, \dots, m\}$, define

$$sp_i(P) = \max\{k < i \mid P[1 \dots k] = P[i - k + 1 \dots i]\},$$

i.e. the length of the longest proper prefix of $P[1 \dots i]$ that is also a suffix of $P[1 \dots i]$.

The key observation is that $sp_i(P)$ can be computed recursively from $sp_{i-1}(P)$, giving an efficient $\mathcal{O}(m)$ preprocessing algorithm. There are two cases:

1. **Extension is possible.** Let $k = sp_{i-1}(P)$. If $P[k+1] = P[i]$, then the longest prefix-suffix of $P[1 \dots i]$ is simply one character longer than that of $P[1 \dots i-1]$, so $sp_i(P) = k + 1$.

Example: with $P = \text{abcdeabdef}$ at $i = 7$, we have $sp_6(P) = 1$ (the prefix-suffix of abcdea is a , length 1). We check $P[2] = \text{b}$ against $P[7] = \text{b}$: they match, so $sp_7(P) = 2$ (see Figure 1.4).

2. **Extension is not possible.** If $P[k+1] \neq P[i]$, we cannot extend the current prefix-suffix. We then ask: what is the longest prefix-suffix of the prefix-suffix itself? In other words, we fall back to $k \leftarrow sp_k(P)$ and try to extend that instead. We repeat this process until either a match is found or we reach $k = 0$, in which case $sp_i(P) = 0$.

Example: still with $P = \text{abcdeabdef}$ at $i = 8$, we have $sp_7(P) = 2$ (the prefix-suffix of abcdeab is ab). We check $P[3] = \text{c}$ against $P[8] = \text{d}$: no match. We fall back to $k \leftarrow sp_2(P) = 0$, so we check $P[1] = \text{a}$ against $P[8] = \text{d}$, still no match. We have reached $k = 0$, so $sp_8(P) = 0$.




Figure 1.4: Computing $sp_7(P)$ for $P = \text{abcdeabdef}$. We know $k = sp_6(P) = 1$: the suffix of $P[1 \dots 6]$ matching a prefix of P has length 1 (namely $\alpha = \text{a}$, violet). We compare $P[k+1] = P[2] = \text{b}$ against $P[i] = P[7] = \text{b}$ (orange): they match, so $sp_7(P) = 2$.

The resulting algorithm is the following:

Algorithm 2: Computing sp values

Input: Pattern $P[1 \dots m]$
Output: Array $sp[1 \dots m]$ where $sp[i] = sp_i(P)$

```

1  $sp[1] \leftarrow 0$ 
2  $k \leftarrow 0$            // length of current candidate prefix-suffix
3 for  $i \leftarrow 2$  to  $m$  do
4   while  $k > 0$  and  $P[k+1] \neq P[i]$  do
5      $k \leftarrow sp[k]$            // fall back
6   if  $P[k+1] = P[i]$  then
7      $k \leftarrow k+1$ 
8    $sp[i] \leftarrow k$ 
9 return  $sp$ 

```

Remark 1.1.8 (Why fall back to $sp[k]$?). At every point in the loop, k has a precise meaning: it is the length of the longest prefix of P that currently matches a suffix of $P[1 \dots i-1]$. In other words,

$$P[1 \dots k] = P[i-k \dots i-1],$$

the k characters immediately before position i are exactly the first k characters of the pattern. When $P[k+1] = P[i]$ we can extend this match by one: $sp[i] = k+1$. When $P[k+1] \neq P[i]$ the match of length k cannot be extended, and we need a shorter candidate $j < k$ that still satisfies the same property: $P[1 \dots j] = P[i-j \dots i-1]$.

Key insight. $P[1 \dots k]$ is not an arbitrary string — it is a *prefix* of P . Any valid candidate j must be a prefix of P that also appears as a suffix of $P[1 \dots i-1]$. Two simple observations handle this:

- *Prefix of a prefix is a prefix:* if $P[1 \dots j]$ is a prefix of $P[1 \dots k]$, it is a prefix of P .
- *Suffix of a suffix is a suffix:* since $P[1 \dots k]$ is a suffix of $P[1 \dots i-1]$ (by the invariant), any suffix of $P[1 \dots k]$ of length j is also a suffix of $P[1 \dots i-1]$.

Therefore j is a valid candidate *if and only if* $P[1 \dots j]$ is a *border* of $P[1 \dots k]$ (a string that is both a prefix and a suffix of it). The longest such border is $sp[k]$, already computed during preprocessing. Any value strictly between $sp[k]$ and k is, by definition of $sp[k]$, not a border of $P[1 \dots k]$ and therefore not a valid candidate — no check is needed. We jump directly to $k \leftarrow sp[k]$ and repeat. Figure 1.5 shows this chain for $i = 6$, $k = 3$.

Note 1.1.9 (A fallback that succeeds). Let $P = \text{ababcababa}$ and consider $i = 10$. We have $k = 4$ since $P[1 \dots 4] = P[6 \dots 9] = \text{abab}$. Check $P[5] = \text{c}$ vs. $P[10] = \text{a}$: mismatch. Could $k' = 3$ be a valid candidate? It would need $P[1 \dots 3]$ to be a border of abab , i.e. $\text{aba} = \text{bab}$: no. The longest border of abab is ab (length 2), so $sp[4] = 2$. We fall back to $k = 2$: $P[1 \dots 2] = P[8 \dots 9] = \text{ab}$. Check $P[3] = \text{a}$ vs. $P[10] = \text{a}$: **match!** Hence $sp[10] = 3$; skipping $k' = 3$ was the only correct move.

FORMAL DETAILS — Why no candidate between $sp[k]$ and k exists

Suppose for contradiction that $sp[k] < k' < k$ and that k' satisfies the same property as k , i.e. $P[1 \dots k'] = P[i - k' \dots i - 1]$. Since $k' < k$, the string $P[i - k' \dots i - 1]$ is a suffix of $P[i - k \dots i - 1] = P[1 \dots k]$. Combined with $P[1 \dots k']$ being a prefix of P , this makes $P[1 \dots k']$ a border of $P[1 \dots k]$ of length $k' > sp[k]$, contradicting $sp[k]$ being the *longest* proper border of $P[1 \dots k]$. $\square$

FORMAL DETAILS — Correctness of the sp preprocessing algorithm

We prove that at the end of iteration i , $sp[i]$ equals the length of the longest proper border of $P[1 \dots i]$.

Precondition $P[1 \dots m]$ is given; $sp[1] = 0$ and $k = 0$ are set before the loop.

Loop guard i ranges from 2 to m .

Invariant At the *start* of iteration i : $k = sp[i-1]$, i.e. k is the length of the longest proper border of $P[1 \dots i-1]$ (and in particular $P[1 \dots k] = P[i-k \dots i-1]$).

Postcondition $sp[i]$ is correctly set for all $1 \leq i \leq m$.

Base case Before the first iteration ($i = 2$), $k = 0 = sp[1]$ since $P[1 \dots 1]$ has no proper border. $\checkmark$

Preservation Assume the invariant holds at the start of iteration i , i.e. $k = sp[i-1]$. The inner **while** loop traverses the border chain $k \rightarrow sp[k] \rightarrow sp[sp[k]] \rightarrow \dots$, stopping as soon as $P[k+1] = P[i]$ or $k = 0$. By the border characterisation (a candidate j is valid iff $P[1 \dots j]$ is a border of $P[1 \dots k]$), this finds the longest extendable border. The subsequent **if** either increments k by one or leaves $k = 0$. In both cases $sp[i] \leftarrow k$ is correct, and $k = sp[i]$ is the invariant for iteration $i+1$. $\checkmark$

Postcondition When the loop exits ($i = m + 1$), all $sp[i]$ have been correctly computed. $\checkmark$

Step-by-step: Computing sp values

Let's trace the algorithm on $P = \text{ababaca}$.

1. $i = 1$: By definition $sp[1] = 0$. $k \leftarrow 0$.
2. $i = 2$: $k = 0$, $P[1] = \text{a} \neq P[2] = \text{b}$. $sp[2] = 0$.
3. $i = 3$: $k = 0$, $P[1] = \text{a} = P[3] = \text{a}$. $k \leftarrow k + 1 = 1$. $sp[3] = 1$.
4. $i = 4$: $k = 1$, $P[2] = \text{b} = P[4] = \text{b}$. $k \leftarrow k + 1 = 2$. $sp[4] = 2$.
5. $i = 5$: $k = 2$, $P[3] = \text{a} = P[5] = \text{a}$. $k \leftarrow k + 1 = 3$. $sp[5] = 3$.
6. $i = 6$: $k = 3$. The invariant gives $P[1 \dots 3] = P[3 \dots 5] = \text{aba}$. We check $P[4] = \text{b}$ against $P[6] = \text{c}$: mismatch, so we fall back. The fallback chain is shown in Figure 1.5:
 - $k \leftarrow sp[3] = 1$: the longest border of **aba** is **a** (length 1). Check $P[2] = \text{b}$ against $P[6] = \text{c}$: mismatch.

- $k \leftarrow sp[1] = 0$: the longest border of **a** is empty. Check $P[1] = \mathbf{a}$ against $P[6] = \mathbf{c}$: mismatch.
- $k = 0$: no candidate left. $sp[6] = 0$.




Figure 1.5: Fallback chain at step $i = 6$ for $P = \mathbf{ababaca}$. The three rows show *all* borders of $P[1..5] = \mathbf{ababa}$ (lengths 3, 1, 0), tried longest first. For each candidate k , the violet cells show the matched suffix $P[i - k..i - 1] = P[1..k]$; the orange cell is the character $P[k + 1]$ we would need at position i to extend it. All three fail, so $sp[6] = 0$.

7. $i = 7$: $k = 0$, $P[1] = \mathbf{a} = P[7] = \mathbf{a}$. $k \leftarrow k + 1 = 1$. $sp[7] = 1$.

i	1	2	3	4	5	6	7
$P[i]$	a	b	a	b	a	c	a
$sp[i]$	0	0	1	2	3	0	1

Remark 1.1.10. The “fall back to the sp value of the current sp value” is the recursive heart of KMP: the same shift logic used during the search phase is also used during preprocessing.

Exercise 1.1.11. Design an algorithm that, given a pattern P of length m , computes $sp_i(P)$ for every $i \in \{1, \dots, m\}$ in $\mathcal{O}(m)$ time.

Hints:

1. Try to express $sp_i(P)$ in terms of $sp_{i-1}(P)$. What is the relationship between the longest prefix-suffix of $P[1 \dots i]$ and that of $P[1 \dots i - 1]$?
2. There are two cases depending on whether $P[sp_{i-1}(P) + 1]$ equals $P[i]$ or not. What can you do in each case?
3. For the harder case, think recursively: if you cannot extend a prefix-suffix of length k , which shorter candidate should you try next?

Exercise 1.1.12 (Period of a string). The *period* of a string S of length n is the smallest integer $p \geq 1$ such that $S[i] = S[i + p]$ for all $1 \leq i \leq n - p$. Equivalently, p is the length of the shortest string Π such that S is a prefix of $\Pi^\infty = \Pi\Pi\Pi\Pi \dots$.

Design an $\mathcal{O}(n)$ -time algorithm that, given S , computes its period.

Hints:

1. Compute the sp values of S .
2. What is the relationship between $sp_n(S)$ and the period of S ? In particular, is $n - sp_n(S)$ always a period?
3. Is it always the *shortest* period? What additional condition on $sp_n(S)$ guarantees minimality?

INTERMEZZO — Amortised analysis — the credit method

Some algorithms contain operations that look expensive in isolation but are cheap *on average* over any sequence of calls. *Amortised analysis* makes this precise by distributing the cost of occasional expensive operations across the many cheap ones that made them possible.

A clean way to formalise this is the **credit method**.

Definition 1.1.13 (Credit method). Assign each operation an *amortised cost* $\hat{c}$ (which we choose freely). When an operation with real cost c runs:

- if $\hat{c} > c$, the surplus $\hat{c} - c$ is deposited as *credits* into a bank;
- if $\hat{c} < c$, the deficit $c - \hat{c}$ is withdrawn from the bank.

Invariant: the bank balance must never go negative.

Consequence: since the bank starts at 0 and never goes below 0,

$$\sum_i c_i \leq \sum_i \hat{c}_i,$$

i.e. the total real cost is bounded by the total amortised cost.

Example: stack with push and multi-pop. Consider a stack supporting two operations:

- PUSH(x): push element x onto the stack; real cost 1.
- MULTI-POP(k): pop the top $\min(k, |S|)$ elements; real cost = number of elements actually popped.

A single MULTI-POP can cost $\Theta(n)$, so a naive worst-case analysis gives $\mathcal{O}(n^2)$ for n operations. Amortised analysis does much better.

Credit assignment. Attach 1 credit to each element at the moment it is pushed:

- PUSH: amortised cost $\hat{c} = 2$ (pay 1 for the real push, deposit 1 credit *on the element*).
- MULTI-POP(k): amortised cost $\hat{c} = 0$ (use the 1 credit stored on each popped element to pay for its removal).

The bank balance equals the number of elements on the stack, which is always ≥ 0 , so the invariant holds.

Example 1.1.14 (Credit trace for a stack). The table below traces the sequence PUSH(a), PUSH(b), PUSH(c), MULTI-POP(2), PUSH(d), PUSH(e), MULTI-POP(3).

Operation	Stack after	Real cost	Amortised cost	Bank
PUSH(a)	[a]	1	2	1
PUSH(b)	[a, b]	1	2	2
PUSH(c)	[a, b, c]	1	2	3
MULTI-POP(2)	[a]	2	0	1
PUSH(d)	[a, d]	1	2	2
PUSH(e)	[a, d, e]	1	2	3
MULTI-POP(3)	[]	3	0	0
Total		10	10	

The bank never goes negative, confirming the invariant. The total real cost 10 is indeed $\leq$ the total amortised cost 10.

Lemma 1.1.15. *Any sequence of n PUSH and MULTI-POP operations on an initially empty stack costs $\mathcal{O}(n)$ in total.*

Proof. There are at most n pushes, each with amortised cost 2, and all MULTI-POPs have amortised cost 0. The total amortised cost is therefore at most $2n$. By the credit-method bound, the total real cost is at most $2n = \mathcal{O}(n)$. $\square$

Application to KMP preprocessing

The sp -computation algorithm runs in $\mathcal{O}(m)$ total time despite the nested loops. We apply the credit method with k as the potential: credits measure how much k can still decrease.

- k increases by at most 1 per outer iteration (there are $m - 1$ iterations, so k grows at most $m - 1$ times in total).
- Each execution of the inner **while** body sets $k \leftarrow sp[k] < k$, strictly decreasing k .
- $k \geq 0$ always holds.

Assign amortised cost 2 to each outer step that increments k (real cost 1 plus deposit 1 credit), and amortised cost 0 to each inner iteration that decrements k (use the saved credit). The bank balance equals $k \geq 0$, so the invariant holds, and the total number of inner iterations is at most the total number of increments, which is $\mathcal{O}(m)$. Together with the $\mathcal{O}(m)$ overhead of the outer loop, the total preprocessing time is $\mathcal{O}(m)$.

Example 1.1.16. Consider the pattern $P = \text{aaaaa}$ of length 5. The sp values are $sp[1] = 0$, $sp[2] = 1$, $sp[3] = 2$, $sp[4] = 3$, $sp[5] = 4$. The variable

k starts at 0 and increases to 1, then to 2, then to 3, then to 4, as we process each character of P . There are no decreases in this case, so the inner loop runs only once per outer iteration, for a total of $\mathcal{O}(m)$ time.

The search phase

Once the sp array is computed, we scan T left-to-right, maintaining j : the number of characters of P currently matched.

Algorithm 3: KMP search

Input: Text $T[1 \dots n]$, pattern $P[1 \dots m]$, precomputed $sp[1 \dots m]$

Output: All starting positions s where P occurs in T

```

1  $j \leftarrow 0$                                 // characters of  $P$  matched so far
2 for  $i \leftarrow 1$  to  $n$  do
3   while  $j > 0$  and  $P[j+1] \neq T[i]$  do
4      $j \leftarrow sp[j]$                     // fall back along the border chain
5   if  $P[j+1] = T[i]$  then
6      $j \leftarrow j+1$ 
7   if  $j = m$  then
8     output  $i - m + 1$                     // occurrence starting here
9      $j \leftarrow sp[j]$                     // continue searching

```

Figure 1.6 traces the search on a concrete example, showing how sp values let the algorithm carry previously matched characters across fallbacks instead of starting from scratch.




Figure 1.6: KMP search for $P = \text{ababc}$ ($m=5$) in $T = \text{ababababc}$ ($n=9$). The sp array for P is shown at the top (cf. the preprocessing in §1.1.3). **Green:** characters newly matched in this phase. **Teal:** characters carried for free from the previous phase via $j \leftarrow sp[j]$. **Red:** mismatch. After each fallback, the algorithm reuses the last $sp[j]$ matched characters without any comparison, then continues from where it left off in T .

The same amortised argument applies: j increases by at most 1 per outer step and each inner fallback strictly decreases $j \geq 0$, so the inner loop runs $\mathcal{O}(n)$ times in total.

Theorem 1.1.17. *The KMP algorithm finds all occurrences of P in T in $\mathcal{O}(n + m)$ time.*

1.1.4 Z-Algorithm

Definition

The Z-algorithm is a second linear-time preprocessing approach. Where $\text{sp}_i(P)$ looks *backward* (longest prefix that is also a suffix of $P[1 \dots i]$), the Z-value Z_i looks *forward* from position i .

Definition 1.1.18 (Z values). Let A be a string of length n . For $i \geq 2$, define

$$Z_i(A) = \begin{cases} 0 & \text{if } A[i] \neq A[1], \\ \max\{\ell > 0 \mid A[1, \ell] = A[i, i + \ell - 1]\} & \text{otherwise.} \end{cases}$$

In words, $Z_i(A)$ is the length of the longest prefix of A that matches A starting at position i , as shown in Figure 1.7.




Figure 1.7: Z values for $A = \text{aabaab}$ ($n = 6$). The prefix $A[1, 3] = \text{aab}$ (violet) reappears at position 4 (teal), giving $Z_4 = 3$. Also $Z_2 = Z_5 = 1$ since $A[1] = A[2] = A[5] = \text{a}$, and $Z_3 = Z_6 = 0$ since $A[3] = A[6] = \text{b} \neq A[1]$. Z_1 is left undefined.

Connection to pattern matching. Choose a separator $\$ \notin \Sigma$ and form $A = P\$T$ (of length $m + 1 + n$). For any index i in the T -part of A (i.e. $i > m + 1$):

$$Z_i(P\$T) = m \iff P \text{ occurs in } T \text{ starting at position } i - m - 1.$$

The separator prevents Z_i from exceeding m , so equality with m detects all occurrences.

FORMAL DETAILS — Proof of the connection

Write $A = P\$T$ and fix $i > m + 1$ (i.e. i is in the T -part of A).

The separator bounds $Z_i \leq m$. If $Z_i > m$ then $A[i \dots i + m] = A[1 \dots m + 1]$, so in particular $A[i + m] = A[m + 1] = \$$. But $i + m > 2m + 1 > m + 1$, so $A[i + m]$ lies in the T -part of A , hence $A[i + m] \in \Sigma$. Since $\$ \notin \Sigma$ we have a contradiction: $Z_i \leq m$.

($\Rightarrow$) Suppose $Z_i = m$. Then $A[i \dots i + m - 1] = A[1 \dots m] = P$. Since $A[i + j - 1] = T[i - m - 1 + j - 1]$ for $j = 1, \dots, m$, this gives

$T[i - m - 1 \dots i - 1] = P$, i.e. P occurs at position $i - m - 1$ in T .

($\Leftarrow$) Suppose P occurs at position $j = i - m - 1$ in T , i.e. $T[j \dots j + m - 1] = P$. Then $A[i \dots i + m - 1] = T[j \dots j + m - 1] = P = A[1 \dots m]$, so $Z_i \geq m$. Combined with $Z_i \leq m$ we get $Z_i = m$. $\square$

Example 1.1.19. Let $P = \text{aba}$ ($m = 3$) and $T = \text{ababac}$ ($n = 6$). Form $A = \text{aba\$ababac}$ and compute Z values for positions $i > 4$:

i	1	2	3	4	5	6	7	8	9	10
$A[i]$	a	b	a	\$	a	b	a	b	a	c
Z_i	—	0	1	0	3	0	3	0	1	0

The two positions with $Z_i = 3 = m$ are $i = 5$ and $i = 7$, corresponding to occurrences of P in T starting at positions $5 - 3 - 1 = 1$ and $7 - 3 - 1 = 3$. Indeed $T[1 \dots 3] = \text{aba}$ and $T[3 \dots 5] = \text{aba}$. $\checkmark$

Computing Z values: the Z -box algorithm

A naive computation of all Z values costs $\mathcal{O}(n^2)$: for each i , compare character by character until a mismatch.

Intuition: reuse past work. Suppose we have processed positions $2, \dots, i-1$ and at some earlier position l we found $Z_l = z$, i.e. $A[l \dots l+z-1] = A[1 \dots z]$. Call the interval

$$[l, r], \quad r = l + z - 1$$

a *Z-box*: a contiguous block that we *know* is an exact copy of the prefix $A[1 \dots z]$.

Now suppose we want Z_i and i falls *inside* the box ($l \leq i \leq r$). Let $k = i - l + 1$ be the *mirror* position of i in the prefix. Since the box says $A[l \dots r] = A[1 \dots r - l + 1]$, we immediately know


$$A[i] = A[k], \quad A[i+1] = A[k+1], \quad \dots, \quad A[r] = A[r-l+1].$$

Whatever is already recorded in Z_k — the length of the match starting at k in the prefix — is available at i *for free*, without any character comparisons. The only open question is whether the match at i extends beyond r , which requires at most one fresh comparison per step.

The Z -box invariant. The algorithm maintains the Z -box $[l, r]$ with the **largest right endpoint** seen so far:

$$r = l + Z_l(A) - 1 \quad (\text{maximised over all processed positions}).$$

Figure 1.8 shows this setup schematically.



At each step i (from 2 to n) exactly one of three cases applies (illustrated in Figure 1.8):

1. $i > r$ (**outside the box**). No cached information is available. Compare $A[1], A[2], \dots$ against $A[i], A[i+1], \dots$ directly until a mismatch; set Z_i to the number of matches. If $Z_i > 0$, update $l \leftarrow i$, $r \leftarrow i + Z_i - 1$.
2. $i \leq r$ (**inside the box**) — let $k = i - l + 1$.
 - (a) $Z_k < r - i + 1$ (**mirror strictly inside**). The Z-match at k ends before the box boundary, and the same mismatch is reflected at i : set $Z_i = Z_k$, no comparisons needed, no update to $[l, r]$.
 - (b) $Z_k \geq r - i + 1$ (**mirror reaches or exceeds the boundary**). We get $Z_i \geq r - i + 1$ for free; continue comparing from $A[r+1]$ until a mismatch, then update $l \leftarrow i$, $r \leftarrow i + Z_i - 1$.

Figure 1.9 traces the *complete* execution on $A = \text{aabaaab}$, showing every step $i = 2, \dots, 7$.




Figure 1.9: Complete Z-box trace on $A = \text{aabaaab}$, steps $i = 2, \dots, 7$. **Teal**: current Z-box $[l, r]$. **Orange**: current position i (teal fill when i is inside the box). **Violet**: mirror $k = i - l + 1$. **Red**: immediate mismatch ($Z=0$). Bottom row: final Z array (non-zero entries highlighted).

Algorithm 4: Z-Algorithm

Input: String $A[1 \dots n]$

Output: Array $Z[2 \dots n]$ with $Z[i] = Z_i(A)$

```

1  $l \leftarrow 0$ 
2  $r \leftarrow 0$ 
3 for  $i \leftarrow 2$  to  $n$  do
4   if  $i > r$  then
5      $Z[i] \leftarrow 0$ 
6     while  $i + Z[i] \leq n$  and  $A[1 + Z[i]] = A[i + Z[i]]$  do
7        $Z[i] \leftarrow Z[i] + 1$ 
8   else
9      $k \leftarrow i - l + 1$ 
10    if  $Z[k] < r - i + 1$  then
11       $Z[i] \leftarrow Z[k]$  // fully inside - no comparisons
12    else
13       $Z[i] \leftarrow r - i + 1$ 
14      while  $i + Z[i] \leq n$  and  $A[1 + Z[i]] = A[i + Z[i]]$  do
15         $Z[i] \leftarrow Z[i] + 1$ 
16  if  $Z[i] > 0$  and  $i + Z[i] - 1 > r$  then
17     $l \leftarrow i$ 
18     $r \leftarrow i + Z[i] - 1$ 
19 return  $Z$ 

```

FORMAL DETAILS — Correctness of the Z-box algorithm

We prove that the algorithm sets $Z[i] = Z_i(A)$ for every $i \geq 2$.

<i>Precondition</i>	$A[1 \dots n]$ is given; $l = r = 0$ before the loop.
<i>Loop guard</i>	i ranges from 2 to n .
<i>Invariant</i>	At the <i>start</i> of iteration i : (i) $Z[j]$ is correctly set for all $2 \leq j < i$, and (ii) $[l, r]$ is the Z-box with the largest right endpoint seen so far, i.e. $A[l \dots r] = A[1 \dots r-l+1]$ (with $l = r = 0$ meaning no box yet).
<i>Postcondition</i>	$Z[i] = Z_i(A)$ for all $i \in \{2, \dots, n\}$.
<i>Base case</i>	Before $i = 2$: Z is empty and $l = r = 0$, so the invariant holds vacuously. ✓
<i>Preservation</i>	Assume the invariant holds at the start of iteration i. We analyse the two cases. <i>Case 1: $i > r$.</i> The while loop compares $A[1], A[2], \dots$ against $A[i], A[i+1], \dots$ directly and stops at the first mismatch, so $Z[i]$ is set to the correct value by definition. If $Z[i] > 0$, the box is updated to $[l, r] = [i, i+Z[i]-1]$, which has a larger r than before. ✓ <i>Case 2a: $i \leq r$ and $Z[k] < r-i+1$ (mirror strictly inside the box).</i> Let $k = i-l+1$. From the invariant, $A[l \dots r] = A[1 \dots r-l+1]$, so $A[i+j] = A[k+j]$ for $j = 0, \dots, r-i$. In particular $A[i \dots i+Z[k]-1] = A[k \dots k+Z[k]-1] = A[1 \dots Z[k]]$. Since $Z[k] < r-i+1$, position $k+Z[k]$ still lies inside $[1, r-l+1]$, so $A[i+Z[k]] = A[k+Z[k]] \neq A[1+Z[k]]$ (the character that ended the match at k). Therefore $Z[i] = Z[k]$. The box $[l, r]$ is not updated since $i+Z[i]-1 < r$. ✓ <i>Case 2b: $i \leq r$ and $Z[k] \geq r-i+1$ (mirror reaches or exceeds the box boundary).</i> The same translation gives $A[i \dots r] = A[1 \dots r-i+1]$, so $Z[i] \geq r-i+1$ is guaranteed. The algorithm initialises $Z[i] = r-i+1$ and extends it by comparing $A[r+1], A[r+2], \dots$ directly until a mismatch, exactly as in Case 1. The box is then updated to $[i, i+Z[i]-1]$. ✓
<i>Postcondition</i>	When the loop exits, all $Z[i]$ have been correctly computed by the case analysis above. ✓

Amortised $\mathcal{O}(n)$ complexity. The variable r is non-decreasing and bounded by n , so it increments at most n times. Every comparison in the **while** loops either advances r (paid by the budget) or ends the step without advancing r (at most one per outer step i). Hence the total number of comparisons is $\mathcal{O}(n)$.

From Z values to sp values

The sp values of P can be read off directly from its Z-array.

Observation 1.1.20. $Z_j(P) = k$ means $P[j \dots j + k - 1] = P[1 \dots k]$, so $P[1 \dots k]$ is a prefix-suffix of $P[1 \dots i]$ for $i = j + k - 1$. We say j maps to i when $i = j + Z_j(P) - 1$.

Proposition 1.1.21. For each $i \in \{2, \dots, m\}$,

$$\text{sp}_i(P) = Z_{j^*}(P), \quad \text{where } j^* = \min\{j \geq 2 \mid j + Z_j(P) - 1 = i\}.$$

Proof. Any j mapping to i gives a prefix-suffix of $P[1 \dots i]$ of length $Z_j(P) = i - j + 1$. The minimum such j maximises $i - j + 1$, yielding the longest prefix-suffix, which is $\text{sp}_i(P)$ by definition. $\square$

Exercise 1.1.22. Using the Z-array of P , design an $\mathcal{O}(m)$ algorithm that computes $\text{sp}_i(P)$ for every $i \in \{1, \dots, m\}$.

1.1.5 Strong borders and real-time KMP

The real-time bottleneck

KMP with plain sp values is *blind* (it never re-reads T) but not *real-time*: a single new character of T can trigger a fallback chain costing $\mathcal{O}(m)$ in the worst case. The fix is to collapse the entire chain into a single precomputed lookup. We achieve this in two logical steps: first by skipping strictly useless borders, and then by precomputing transitions for every possible character.

Strong border (sp'): Skipping useless fallbacks

Definition 1.1.23 (Strong border). For $i \in \{1, \dots, m - 1\}$, define

$$\text{sp}_i(P)' = \max\{k \leq \text{sp}_i(P) \mid k = 0 \text{ or } P[k + 1] \neq P[i + 1]\}.$$

This is the length of the longest proper prefix-suffix of $P[1 \dots i]$ whose next character differs from $P[i + 1]$.

In plain words, $\text{sp}_i(P)'$ is the length of the longest border of $P[1 \dots i]$ that would not cause an immediate mismatch at $P[i + 1]$.

Any mismatch on $P[i + 1]$ with a text character x would also fail at the next border $k = \text{sp}_i(P)$ if $P[k + 1] = P[i + 1]$ — we say such borders are *useless* (see Figure 1.10). What is a useless border? Example 1.10 shows $P = \text{aabaa}$ at $i = 4$: the border $\alpha = \text{a}$ of length $k = 1$ is useless since $P[k + 1] = P[2] = \text{a}$ equals $P[i + 1] = P[5] = \text{a}$. If a mismatch occurs at $P[5]$ with some text character $x \neq \text{a}$, then after falling back to the border α , we would immediately fail again at $P[k + 1]$ — the border is useless. sp' allows us to skip an entire chain of useless borders in a single step, jumping directly to the longest safe border that could potentially match the next text character.




Figure 1.10: Strong border for $P = \text{aabaa}$ at $i = 4$. The border $\alpha = \text{a}$ (violet, $k = \text{sp}_4(P) = 1$) satisfies $P[1] = P[4]$, but its “next character” $P[k+1] = P[2] = \text{a}$ (orange) equals $P[i+1] = P[5] = \text{a}$. Any mismatch on $P[5]$ with text character $x \neq \text{a}$ would *also* fail immediately at $P[k+1]$ after the fallback — the border is useless. Hence $\text{sp}_4(P)' = 0$: skip straight to the start.

$\text{sp}_i(P)'$ skips the entire chain of useless borders in one step, but it only encodes a single fallback state — the longest “safe” border for the specific character $P[i+1]$. For a completely arbitrary text character x , a fallback chain might still occur. To achieve strict $\mathcal{O}(1)$ time per character, a full character-indexed table is needed.

Character-indexed borders and the DFA: True Real-Time

For a complete real-time solution we precompute, for each position i and character $x \in \Sigma$, the longest border of $P[1 \dots i]$ whose next character is exactly x :

Definition 1.1.24 ($\text{sp}_{i,x}$ values). For $i \in \{0, 1, \dots, m\}$ and $x \in \Sigma$,

$$\text{sp}_{i,x}(P) = \max\{k \leq \text{sp}_i(P) \mid P[k+1] = x\},$$

with default 0 when no such k exists.

This generalises $\text{sp}_i(P)'$: where the strong border encodes a single fallback state for the character $P[i+1]$, the table $\text{sp}_{i,x}$ encodes the optimal fallback for *every* $x \in \Sigma$ simultaneously.

Example 1.1.25 (sp' vs. $\text{sp}_{i,x}$ for $P = \text{aabaab}$). Consider $P = \text{aabaab}$ ($m = 6$) over $\Sigma = \{\text{a}, \text{b}, \text{c}\}$. The sp and sp' values are:

i	1	2	3	4	5	6
$P[i]$	a	a	b	a	a	b
$\text{sp}_i(P)$	0	1	0	1	2	3
$\text{sp}_i(P)'$	0	1	0	0	1	–

At $i = 5$ the border aa ($\text{sp}_5(P) = 2$) is useless because $P[3] = \text{b} = P[6]$; the shorter border a ($k = 1$) satisfies $P[2] = \text{a} \neq P[6]$, so $\text{sp}_5(P)' = 1$.

Now suppose we have matched $P[1..5]$ (state $j = 5$) and the next text character is x . With $\text{sp}_5(P)' = 1$, KMP falls back to state 1 for *every* mismatch $x \neq P[6] = \text{b}$, and may still need further comparisons depending on x :

At $i = 5$, the border chain of $P[1..5] = \text{aabaab}$ is $k \in \{2, 1, 0\}$. Both $k = 2$ ($P[3] = \text{b} = P[6]$) and $k = 1$ ($P[2] = \text{a}$) must be inspected; $\text{sp}_5(P)' = 1$ skips past the useless $k = 2$.

$sp_{5,a} = 1$ because the border of length 1 satisfies $P[2] = a = x$; no border has $P[k+1] = c$, so $sp_{5,c} = 0$.

- $x = a$: $5 \rightarrow 1$, then $P[2] = a = x \Rightarrow$ state 2. **Two comparisons.**
- $x = c$: $5 \rightarrow 1$, $P[2] \neq c$; $1 \rightarrow 0$, $P[1] \neq c$; stay at 0. **Three comparisons.**

By contrast, $sp_{5,x}$ resolves every case in one lookup:

x	a	b	c
$sp_{5,x}$	1	— (match $P[6]$)	0
DFA $\delta(5, x)$	2	6 (full match!)	0

$\delta(5, a) = 2$ replaces two comparisons with one table read; $\delta(5, c) = 0$ replaces three.

The DFA has $m + 1$ states and $|\Sigma|$ transitions per state, so the table takes $\mathcal{O}(m|\Sigma|)$ space and time to build.

The table $\{sp_{i,x}(P)\}_{i,x}$ is precisely the *transition function* of the deterministic finite automaton (DFA) for pattern P : state i encodes “ i characters of P matched so far”, and on reading x the automaton transitions to $sp_{i,x}(P) + \mathbf{1}[P[i+1] = x]$. With this table, each character of T is processed in strict $\mathcal{O}(1)$ worst-case time, making KMP truly **real-time**.

1.1.6 A Comprehensive Trace: KMP and Z-Algorithm

To solidify these concepts, let’s trace KMP and then contrast it directly with the Z-Algorithm.

KMP Step-by-Step

Consider the pattern $P = \text{ababa}$ and the text $T = \text{ababcababa}$. Here we trace the standard KMP behavior (using plain sp values to illustrate the fallback mechanism that DFA optimizes).

For $i = 4$, $P[1..4] = \text{abab}$. Proper prefix-suffixes are ab (len 2) and a (len 1). Max is 2.

1. Preprocessing P . We compute the sp array for P .

i	1	2	3	4	5
$P[i]$	a	b	a	b	a
$sp_i(P)$	0	0	1	2	3

2. Searching T . We maintain a pointer i into T and a pointer j representing the length of the current match in P . The complete trace is shown in Figure 1.11.




Figure 1.11: KMP trace for $P = \text{ababa}$ in $T = \text{ababcbabab}$. Step 1 finds a mismatch at $T[5]$. KMP falls back twice ($j = 4 \rightarrow 2 \rightarrow 0$) without re-reading T . Finally, Step 4 finds a complete match starting at $T[6]$.

Online vs. Real-Time: Bridging the Gap

The KMP trace above highlights a crucial property: KMP processes T one character at a time without ever stepping backwards in T . This makes KMP an **online** algorithm. However, as seen in Steps 1-3, processing the single character $T[5] = c$ required multiple internal operations (j falling back multiple times). If we had used the full DFA ($sp_{i,x}$) discussed previously, those steps would collapse into $\mathcal{O}(1)$ worst-case time, upgrading the algorithm from merely online to strictly **real-time**.

By contrast, the Z-algorithm is *neither* real-time nor strictly online in the same way. The computation of the Z-box $[l, r]$ depends on reading characters ahead in the string. Let's trace it to see this structural difference in action.

Z-Algorithm Trace

Let's compute the Z-values for the string $S = \text{aababca}$. Notice how the algorithm looks ahead to establish the rightmost boundary r .

- $i = 2$: $S[2] = a = S[1]$. Check $S[3] = b \neq S[2]$. $Z_2 = 1$. Box $[l, r] = [2, 2]$.
- $i = 3$: $i > r = 2$. $S[3] = b \neq S[1]$. $Z_3 = 0$. Box stays $[2, 2]$.
- $i = 4$: $i > r = 2$. $S[4] = a = S[1]$. Check $S[5] = b = S[2]$, $S[6] = c \neq S[3]$. $Z_4 = 2$. New box $[l, r] = [4, 5]$.
- $i = 5$: $i \leq r = 5$. Mirror $k = i - l + 1 = 5 - 4 + 1 = 2$. $Z_k = Z_2 = 1$. Remaining box width $r - i + 1 = 5 - 5 + 1 = 1$. Since $Z_k \not\leq r - i + 1$ (Case 2b), we try to extend from $r + 1 = 6$. $S[6] = c \neq S[3]$. No extension. $Z_5 = 1 + 0 = 1$. Box stays $[4, 5]$.
- $i = 6$: $i > r = 5$. $S[6] = c \neq S[1]$. $Z_6 = 0$.
- $i = 7$: $i > r = 5$. $S[7] = a = S[1]$. No more chars. $Z_7 = 1$.

i	1	2	3	4	5	6	7
$S[i]$	a	a	b	a	b	c	a
Z_i	-	1	0	2	1	0	1

1.1.7 Boyer-Moore Algorithm

So far we have seen algorithms that scan the text strictly left to right, never revisiting a character once processed. The Boyer-Moore algorithm takes a different approach: it compares the pattern against the text *from right to left* while still sliding the pattern from left to right. Mismatches on the rightmost characters of the pattern are the cheapest failures possible—they let the algorithm skip large portions of the text.

Intuition

The algorithm combines two independent heuristics. At every alignment it takes the *maximum* shift allowed by either rule:


- **Bad Character Rule (BCR).** When character x in the text causes a mismatch at position i in the pattern, shift P right so that the rightmost occurrence of x in $P[1 \dots i-1]$ aligns with x in the text. If x does not appear in P at all, skip past x entirely.
- **Good Suffix Rule (GSR).** When the suffix $S = P[i \dots m]$ has already been matched and position $i-1$ causes a mismatch, shift P right so that another copy of S in P —preceded by a character *different* from $P[i-1]$ —aligns with the matched copy of S in the text.

Both rules are correct in the sense that neither one ever skips a genuine occurrence of P . At each step Boyer-Moore applies whichever rule yields the larger shift.

Bad Character Rule

Definition 1.1.26 (Bad Character table). For each position $1 \leq i \leq m$ and character $x \in \Sigma$, let $BC(i, x)$ be the largest $j < i$ such that $P[j] = x$, or 0 if no such j exists. When a mismatch occurs at $P[i]$ against text character x , shift P right by $i - BC(i, x)$ positions.

Example 1.1.27 (BCR skip). Search $P = \text{ALGO}$ in $T = \dots \text{ALBERO} \dots$. Aligning P with the first four characters, we compare right-to-left: $P[4] = \text{O}$ vs $T[4] = \text{E}$ is a mismatch. Since E does not appear anywhere in P , BCR shifts the entire pattern past the mismatch character.



Scanning right-to-left is the key insight: a mismatch at the last character of the current alignment gives the most information about how far to jump.

Preprocessing options. Two standard approaches exist for building the BCR table:

- a) **Full 2-D table** ($|\Sigma| \times m$): $\mathcal{O}(1)$ query time, but $\mathcal{O}(|\Sigma|m)$ space.
- b) **Position lists**: for each character store the sorted list of its occurrences in P and answer queries with binary search in $\mathcal{O}(\log m)$.


Remark 1.1.28. Although binary search costs $\mathcal{O}(\log m)$ per query, the large shifts that BCR produces amortise this overhead over the full run of the algorithm, making the two implementations equivalent in practice.

Good Suffix Rule

After matching suffix $S = P[i \dots m]$, we want to shift P so that another copy of S aligns with the text. We require that copy to be preceded by a character *different* from $P[i-1]$; otherwise the same mismatch would recur immediately.

Definition 1.1.29 ($L'(i)$). $L'(i)$ is the largest index $j < m$ such that

1. $P[i \dots m]$ occurs in P ending at position j , and
2. the character immediately before that occurrence differs from $P[i-1]$: $P[j - (m - i)] \neq P[i-1]$.



To compute $L'(i)$ efficiently we exploit its connection with the Z-algorithm applied to the reversed pattern P^r .

Definition 1.1.30 ($N_j(P)$). $N_j(P)$ is the length of the longest suffix of $P[1 \dots j]$ that is also a suffix of P .

Theorem 1.1.31. The values $N_j(P)$ satisfy

$$N_j(P) = Z_{m-j+1}(P^r),$$

where $m = |P|$ and P^r denotes P reversed. Consequently, the entire GSR table can be computed in $\mathcal{O}(m)$ time by running the Z-algorithm on P^r .

FORMAL DETAILS — **Proof:** $N_j(P) = Z_{m-j+1}(P^r)$

Set $k = m - j + 1$ and recall the two definitions:

- $N_j(P) = \ell$ iff ℓ is the largest integer such that $P[j-\ell+1 \dots j] = P[m-\ell+1 \dots m]$, i.e. the last ℓ characters of $P[1 \dots j]$ equal the last ℓ characters of P .
- $Z_k(P^r) = \ell$ iff the length- ℓ prefix of P^r matches P^r at position k :

$$P^r[1 \dots \ell] = P^r[k \dots k+\ell-1].$$

Key observation. Since $P^r[i] = P[m+1-i]$, reading P^r from left to right is the same as reading P from right to left. More precisely, for any $\ell \geq 1$ and $s = 0, 1, \dots, \ell-1$:

$$P^r[1+s] = P[m-s] \quad \text{and} \quad P^r[k+s] = P[m+1-k-s] = P[j-s],$$

where the last equality uses $k = m-j+1$.

The two conditions are equivalent. For any $\ell \geq 0$:

$$\begin{aligned} P^r[1 \dots \ell] = P^r[k \dots k+\ell-1] &\iff P[m-s] = P[j-s] \forall s \in \{0, \dots, \ell-1\} \\ &\iff P[m-\ell+1 \dots m] = P[j-\ell+1 \dots j]. \end{aligned}$$


The rightmost condition is exactly the requirement that $P[1 \dots j]$ has a suffix of length ℓ that matches a suffix of P .

Conclusion. Because the two conditions agree for every ℓ , their maxima agree:

$$N_j(P) = Z_{m-j+1}(P^r). \square$$

Complexity. Running the Z-algorithm on P^r costs $\mathcal{O}(m)$ time and produces all m values $Z_1(P^r), \dots, Z_m(P^r)$ in a single pass. Reading off $N_j = Z_{m-j+1}$ for $j = 1, \dots, m$ is then an $\mathcal{O}(m)$ post-processing step, so the full GSR precomputation is $\mathcal{O}(m)$ overall.

Reusing the Z-algorithm on P^r means no new code is needed for GSR preprocessing: run the same routine on the reversed pattern.



Complexity

Theorem 1.1.32. *In the worst case, Boyer-Moore runs in $\mathcal{O}(nm)$ time.*

The classical worst case is $T = a^n$, $P = ba^{m-1}$: every alignment compares the entire pattern right-to-left, reaches the first character, and shifts by only one position.

Remark 1.1.33. The **Apostolico–Giancarlo** extension records which portions of the text have already been matched, avoiding redundant comparisons and bringing the worst-case complexity down to $\mathcal{O}(n+m)$. On the other hand, in favourable inputs (e.g. $T = c^n$, $P = a^{m-1}c$) BCR allows Boyer-Moore to scan the text in as few as $\lceil n/m \rceil$ character comparisons—sublinear in n .

Chapter 2

Randomized Algorithms

Chapter 3

Data Compression